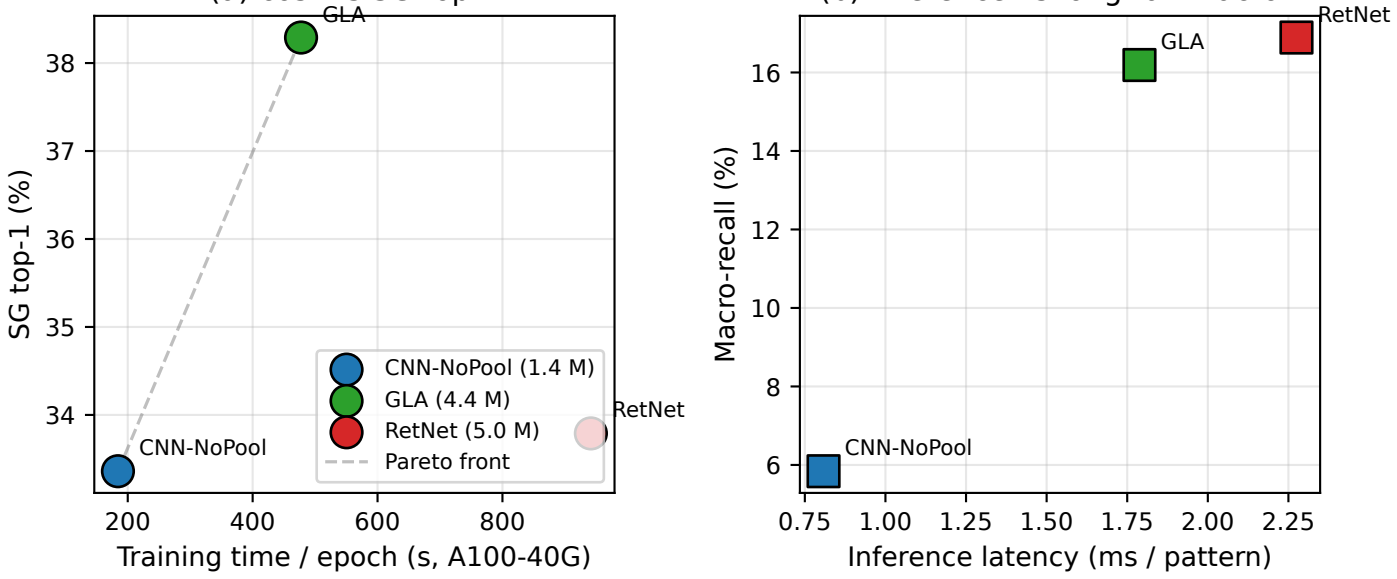
(a) cost vs SG top-1
GLA
CNN-NoPool
RetNet
CNN-NoPool (1.4 M)
GLA (4.4 M)
RetNet (5.0 M)
Pareto front
SG top-1 (%)
Training time / epoch (s, A100-40G)
(b) inference vs long-tail macroR
RetNet
GLA
CNN-NoPool
Macro-recall (%)
Inference latency (ms / pattern)